

Statistical NLP Part-1 - Text Classification

- **DOMAIN:** Digital content management

- **CONTEXT:**

Classification is probably the most popular task that you would deal with in real life. Text in the form of blogs, posts, articles, etc. is written every second. It is a challenge to predict the information about the writer without knowing about him/her. We are going to create a classifier that predicts multiple features of the author of a given text. We have designed it as a Multi label classification problem

- **DATA DESCRIPTION:**

Over 600,000 posts from more than 19 thousand bloggers The Blog Authorship Corpus consists of the collected posts of 19,320 bloggers gathered from blogger.com in August 2004. The corpus incorporates a total of 681,288 posts and over 140 million words - or approximately 35 posts and 7250 words per person. Each blog is presented as a separate file, the name of which indicates a blogger id# and the blogger's self-provided gender, age, industry, and astrological sign. (All are labelled for gender and age but for many, industry and/or sign is marked as unknown.) [\[Source\]](#)

All bloggers included in the corpus fall into one of three age groups:

- 8240 "10s" blogs (ages 13-17)
- 8086 "20s" blogs(ages 23-27) and
- 2994 "30s" blogs (ages 33-47)

For each age group, there is an equal number of male and female bloggers. Each blog in the corpus includes at least 200 occurrences of common English words. All formatting has been stripped with two exceptions. Individual posts within a single blogger are separated by the date of the following post and links within a post are denoted by the label url link.

- **PROJECT OBJECTIVE:**

The need is to build a NLP classifier which can use input text parameters to determine the label/s of the blog.

In [1]:

```
# !conda install dtale -c conda-forge
# if you want to also use "Export to PNG" for charts
# !conda install -c plotly python-kaleido
```

In [2]:

```
# imports

import os
import dtale
import random
import warnings
from time import time
from math import floor
from pathlib import Path
import pandas as pd, numpy as np
from pprint import pprint
import matplotlib.pyplot as plt
import seaborn as sns
from tqdm import tqdm
from collections import defaultdict
import tensorflow as tf
tqdm.pandas()
warnings.filterwarnings('ignore')
%matplotlib inline
```

In [3]:

```
# reproducibility
seed = 7
random.seed(seed)
```

- **Import the data.**

In [4]:

```
blog_df = pd.read_csv('./data/blogtext.csv')
d = dtale.show(blog_df)
d
```

```
2021-07-25 18:25:51,567 - INFO      - Note: NumExpr detected 16 cores but "NUMEXPR_MAX_THR
EADS" not set, so enforcing safe limit of 8.
2021-07-25 18:25:51,568 - INFO      - NumExpr defaulting to 8 threads.
```

Out[4]:

In [5]:

```
# 681284 blog texts
blog_df.shape
```

Out[5]:

```
(681284, 7)
```

In [6]:

```
blog_df.isna().sum()  # no null values, dataset is good in terms of data completeness
```

Out[6]:

```
id          0
gender      0
age         0
topic       0
sign        0
data       0
```

```
date      0
text      0
dtype: int64
```

In [7]:

```
blog_df.info()  # topic, sing columns to be converted to categorical and date can be split into day, month, year if required for analysis
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 681284 entries, 0 to 681283
Data columns (total 7 columns):
#   Column  Non-Null Count  Dtype
---  ------  -
0    id      681284 non-null   int64
1    gender  681284 non-null   object
2    age     681284 non-null   int64
3    topic   681284 non-null   object
4    sign    681284 non-null   object
5    date    681284 non-null   object
6    text    681284 non-null   object
dtypes: int64(2), object(5)
memory usage: 36.4+ MB
```

In [8]:

```
blog_df['id'] = pd.Categorical(blog_df.id)
blog_df['topic'] = pd.Categorical(blog_df.topic)
blog_df['sign'] = pd.Categorical(blog_df.sign)
blog_df['text'] = blog_df['text'].fillna('').apply(str)
```

In [9]:

```
blog_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 681284 entries, 0 to 681283
Data columns (total 7 columns):
#   Column  Non-Null Count  Dtype
---  ------  -
0    id      681284 non-null   category
1    gender  681284 non-null   object
2    age     681284 non-null   int64
3    topic   681284 non-null   category
4    sign    681284 non-null   category
5    date    681284 non-null   object
6    text    681284 non-null   object
dtypes: category(3), int64(1), object(3)
memory usage: 24.0+ MB
```

In [10]:

```
# Imbalanced classes!
blog_df.topic.value_counts()
```

Out[10]:

```
indUnk      251015
Student     153903
Technology   42055
Arts         32449
Education    29633
Communications-Media  20140
Internet     16006
Non-Profit   14700
Engineering  11653
Law          9040
Publishing   7753
Science      7269
Government   6907
Consulting   5862
Religion     5235
Fashion      4251
```

Fashion	4031
Marketing	4769
Advertising	4676
BusinessServices	4500
Banking	4049
Chemicals	3928
Telecommunications	3891
Accounting	3832
Military	3128
Museums-Libraries	3096
Sports-Recreation	3038
HumanResources	3010
RealEstate	2870
Transportation	2326
Manufacturing	2272
Biotech	2234
Tourism	1942
LawEnforcement-Security	1878
Architecture	1638
InvestmentBanking	1292
Automotive	1244
Agriculture	1235
Construction	1093
Environment	592
Maritime	280

Name: topic, dtype: int64

In [11]:

```
blog_df.sign.value_counts()
```

Out[11]:

Cancer	65048
Aries	64979
Taurus	62561
Libra	62363
Virgo	60399
Scorpio	57161
Pisces	54053
Leo	53811
Gemini	51985
Sagittarius	50036
Aquarius	49687
Capricorn	49201

Name: sign, dtype: int64

In [12]:

```
# !pip install langdetect
```

In [13]:

```
# Looking for languages
from langdetect import detect_langs, LangDetectException
from collections import defaultdict

languages = []
lang_samples = defaultdict(list)
# pick random texts out of the dataset as language detection on every text is not feasible
samples = np.random.choice(len(blog_df.text), size=5000, replace=False)
# Loop over the rows of the dataset and append
for row in tqdm(samples):
    try:
        text = blog_df.text[row]
        lang = detect_langs(text)
        clean_lang = str(lang).split(':')[0][1:]
        lang_samples[clean_lang].append(text)
        languages.append(clean_lang)
    except LangDetectException:
        pass
```

```
100%|██████████████████████████████████████████████████████████████████████████████| 5000/  
5000 [00:28<00:00, 174.63it/s]
```

In [14]:

```
en
["      NOTE: You may want to edit your posts so they don't include the HeppSTAW na
me, folks... would you really want (P)ank (Fr)ate finding out how much you weight? And si
nce blogs ARE included in most search engines....      "]
```

```
lt
['      urlLink      &nbsp;    urlLink      ']
```

```
t1
['      Hintayin ko na tanong ni Debb para maging organized ang game.      ']
```

```
af
['      urlLink      Room 309, The Premier Room&nbsp;   urlLink      ']
```

```
da
['      *deleted*      ']
```

```
hu
['      urlLink      Get A Job      ']
```

id
[" leaving dundee tomorrow. yeaah!!! waitin' to see kak minn cos she'll b
e leaving dundee in august for good.. isk isk... my sis angkat. One of the reasons why i
chose dundee over aberdeen and leeds. anyways.. the following has been copied from www.
utusan.com.my. thought i'd paste it on my blog for my own personal reading in future.. ha!
ahahah. MY DAD SHOULD REALLY READ THIS!!! especially the last bit! Doktor pakar diberi
layanan kelas kedua kenaikan pangkat Saudara Pengarang, KENYATAAN Menteri Kesihatan di Pa
rlimen pada 7 Julai 2004 yang disiarkan oleh akhbar Utusan Malaysia adalah dirujuk. Belia
u mendakwa sehingga 31 Mei 2004, masih terdapat 875 jawatan pegawai perubatan pakar pelba
gai bidang yang belum diisi daripada sejumlah 2,397 jawatan pegawai perubatan pakar di Ke
menterian Kesihatan Malaysia. Di sini, saya dengan sebulat suara menyokong kenyataan be
liau dan ia amat berpatutan dengan iklan yang dikeluarkan oleh Kementerian Kesihatan yang
disiarkan di laman web kementerian (untuk sesiapa yang berminat membacanya, [www.moh.gov.m](http://www.moh.gov.my)
[y](http://www.moh.gov.my)). Namun begitu, janganlah rakyat Malaysia terkejut jika semakin ramai pegawai peruba

y, . Namun begitu, janganlah rakyat Malaysia terkejut, jika semakin ramai pegawai perubatan pakar yang akan meletakkan jawatan mereka disebabkan iklan ini. Untuk pengetahuan semua, gelaran doktor perubatan adalah bersamaan dengan pegawai perubatan di dalam skim perkhidmatan kerajaan. Untuk pengetahuan pembaca, sebelum Skim SSM (Sistem Saraan Malaysia) diperkenalkan, kenaikan pangkat untuk jawatan pegawai perubatan pakar di dalam Skim SSB (Sistem Saraan Baru) adalah dari gred U3 ke gred U2. Gaji untuk pegawai perubatan pakar klinikal gred U3 untuk skim SSB adalah di dalam lingkungan RM5,012.81 (iaitu selepas mendapat hadiah dua kenaikan gaji dan bayaran insentif pakar apabila seseorang pegawai perubatan lulus ijazah sarjana kepakaran dan diwartakan sebagai pakar). Apabila dinaikkan pangkat ke gred U2 di dalam SSB, seorang pegawai perubatan pakar gred U2 akan menerima imbuhan gaji di dalam lingkungan RM6,988.51. Kenaikan pangkat ke gred U2 ini akan dinikmati oleh pegawai perubatan pakar tersebut setelah dua hingga tiga tahun diwartakan sebagai pakar klinikal di gred U3. Ini merupakan pertambahan gaji sebanyak RM1,975.70, satu jumlah yang agak berpatutan. Bagaimanapun kami hanya nikmati setelah terpaksa bertungkus lumus untuk menamatkan pengajian di peringkat sarjana kepakaran yang mengambil masa selama empat hingga lima tahun (tidak termasuk ijazah sarjana muda doktor perubatan, M.D, M.B.B.S, yang telah diambil dalam masa enam hingga tujuh tahun). Tetapi lain pula halnya di dalam SSM yang diperkenalkan pada penghujung tahun 2002, jawatan kami sebagai pegawai perubatan pakar klinikal gred U41 hanya dinaikkan ke pegawai perubatan pakar gred U44 (sepatutnya jawatan setara yang perlu ditawarkan untuk kenaikan pangkat di dalam SSM untuk pegawai perubatan pakar adalah pada gred U48 iaitu bersamaan gred U2 untuk skim SSB). Gaji yang kami nikmati untuk jawatan pegawai perubatan pakar klinikal gred U41 adalah di dalam lingkungan RM5,177.81 dan apabila kami dinaikkan pangkat ke gred U44 (selepas diwartakan dan semestinya perlu lulus peperiksaan Penilaian Tahap Kecekapan, SSM), gaji baru yang akan kami nikmati hanyalah RM5,631.23. iaitu kenaikan sebanyak RM453.42. Bayangkanlah pertambahan gaji sebanyak itu yang kami akan dapat di dalam Skim SSM setelah bertungkus lumus belajar untuk mendapatkan ijazah kepakaran. Untuk pengetahuan pembaca, gaji pegawai perubatan pakar gred U48 skim SSM adalah bermula daripada RM7,303.51. Suatu perbezaan yang amat ketara di dalam skim SSM yang diperkenalkan berbanding skim SSB. Pada pendapat saya, jika seseorang hanyalah pegawai perubatan gred U41 yang tidak mempunyai ijazah kepakaran dan dinaikkan ke gred U44, itu merupakan satu rezeki. Tetapi yang saya tidak faham, kenapalah kenaikan pangkat ke gred U44 ini diperkenalkan kepada pegawai perubatan yang mempunyai kepakaran klinikal. Kenapa, pihak kementerian tidak mahu memperkenalkan kembali skim kenaikan pangkat ke gred kanan kepada pegawai perubatan, seperti yang diamalkan sewaktu Jawatankuasa Kabinet 1976, yang memberikan kenaikan pangkat kepada semua pegawai perubatan yang telah berkhidmat selama lima tahun. Gred U43/U44 boleh digunakan sebagai gred kanan. Manakala Gred U47/U48 dan ke atas dijadikan gred kenaikan pangkat untuk pegawai perubatan pakar dan juga kepada pegawai perubatan yang cemerlang dan sebagainya. Sebagai kesimpulan dari tulisan saya ini, 1. Pihak kementerian sememangnya menggalakkan pegawai perubatan pakar rakyat Malaysia untuk berhijrah, kerana mereka boleh mendapatkan khidmat pegawai perubatan pakar dari luar. Yang menariknya di sini untuk khidmat pakar dari luar, kementerian menawarkan Gred permulaan U48 sehingga ke gred U54 untuk pegawai-pegawai perubatan pakar dari luar ini, yang sesetengahnya tidak diketahui asal-usul ijazah kepakaran mereka ambil. Bukanlah pegawai perubatan pakar ini tidak mahu berkhidmat di hospital-hospital kerajaan, tetapi sebenarnya kami merasakan yang kami ini dianak-tirikan dan diberi layanan kelas kedua oleh Kementerian Kesihatan. 2. Hospital-hospital kerajaan akan sentiasa dianggap sebagai hospital kelas kedua oleh rakyat Malaysia. Ini kerana kebanyakan penghijrahan pakar ini merupakan pakar-pakar muda yang bercita-cita tinggi untuk menjadi seorang pakar khusus/sub-kepakaran. Dengan itu, hospital kerajaan akan tinggal dengan pakar-pakar perunding yang akan bersara dan tiada kelebihan di bidang sub-kepakaran berbanding pegawai perubatan pakar yang masih muda. Jikapun mereka ini pergi melanjutkan pengajian di bidang sub-kepakaran, berapa lama sangat yang mereka dapat berkhidmat. Selepas itu mereka pun bersara dan menyertai pihak swasta dengan sub-kepakaran mereka. Inilah fenomena yang kita alami sekarang ini di mana hospital-hospital swasta disanjung tinggi oleh rakyat, walaupun terpaksa membayar kos rawatan yang begitu tinggi. Dari manakah datangnya pakar yang berkhidmat di hospital swasta ini, jika bukan daripada penghijrahan dari sektor awam. 3. Untuk adik-adik yang terlalu amat meminati bidang perubatan. Bidang perubatan bukanlah bidang yang 'glamour'. Ia memerlukan jiwa yang kental dan masa yang amat-amat panjang untuk mengharungi bidang ini. Janganlah mengharap imbuhan yang luar biasa dari bidang ini (kecualilah jika anda berazam untuk meninggalkan sektor awam selepas tamat belajar). Teruskan niat jika anda mencintai bidang perubatan yang sememangnya memerlukan ketabahan mental dan fizikal. - DOKTOR PAKAR KERAJAAN, Batu Ferringhi, Pulau Pinang.

"]

so
[' o man ooo man ']

sk
[' zoom ']

cy
[' hey o.d. ']

```

fr
["          That's just noise, Dave.          "]

es
["          la vida es una cosa curiosa... y ms cuando uno acaba de comer, justo ahorita todo se siente 'lento' y no dan ganas de nada... solo de pensar... probablemente por eso varios filosofos eran 'rellenitos'... se la pasaban comiendo y gozando de ese curioso estado de letargo que te crea la sensacin de estar saciado por una buena comida... lo nico que faltó fue el vino... un buen vino tinto junto con esa lasagna que nos acabamos de hechar hubiera sido la onda... y es que bien dicen 'despus de comer, ni una carta leer'.. he he he... as que yo solo y mis pensamientos nos la pasamos muy bien... no entiendo como todo mundo se puede poner a trabajar luego luego... ESO NO ESTA BIEN!!... debe uno de dejar descansar la comida... por cierto, hablando de comida y otros pecados, mi cinturón que no creo que heche mentiras, ya me dice que definitivamente mi talla esta cambiando.. por ms que la báscula quiera verse condescendiente creo que definitivamente estoy subiendo de peso... muy triste!!... mas que nada por mi espalda y todas esas cosas... el lunes comenzar una especie de dieta, que ms bien va a constar de eliminar de mi hábitos todos esos gansitos y cocas que he consumido ultimaente... me va a costar trabajo dejarlos... pero va a ser necesario...          "]

fi
['          How did you put yours on?  Telllllllllllllllllllllllll meeeeeeeeeeeee.  MOOOOOOOOO OOOOOOORTAAAAAATAAAAAAAR!          ']

pt
['          PARENTAL          ADVISORY          OODREAMS0COME0TRUE0 CONTAINS EXPLICIT LYRICS
Username:          From  urlLink Go-Quiz.com          ']

it
['          Hello          ']

tr
['          urlLink          Smiley happy Delu          ']

ru
['          тп а          ұртғһсoŁkeeeeP  «·ēŁXçζg          ']

ro
['          Straight people  urlLink are so cute ...          ']

ko
['          最終 只要對得起自己 就夠了吧 ... 對於那些 睜眼說瞎話的混蛋 ..... 即使死不承認說過哪些話 作過哪些事 其實他們心裡明明 深深記得的 呼嚨 只是他們保護自己的手段 裝死 只是他們最後防衛的武器 在裝死的瞬間 才能變態的得到一點點解脫 可憐的是他們 活在自以為的世界裡 永遠逃不出來 在對不起良心的地獄裡 死過千遍萬遍          ']

vi
['          Tic Toc          ']

zh-cn
['          [ 亲近则易生悔慢之心 ] 有时候,我们会对别人授予的小惠[ 感激不尽 ], 欲对亲人,父母的一辈子恩情[ 视而不见 ]! 因为,亲人对我们来说,太亲近了,说以我们看不见他们的好、看不见他们的牺牲!我们把亲人辛苦的照顾、爱护,视为[ 理所当然 ]!有时为了小争执,还有[ 憎恨]他们,而不知[ 心存感激 ]! 西方有句话说,[ 上帝又两个住处,一个是在天堂,另一个是在感恩者的心里 !] 是的,常[ 心存感谢 ]的人是快乐的;一颗[ 知恩的心 ],也就是[ 快乐的心 ],上帝将住在其中! 且让我们学习—[ 多想想别人的好,忘记别人的不好 ]、[ 对所得恩典的谢意,能与争取恩典时一样地热心 ],则我们心中将更充满爱,也将如在天堂一样快乐。          ']

sl
['          I need some sake soon.....ZzzzzZzzz          ']

```

- *There are multiple languages and character sets in the dataset!*

• Visualize

In [15]:

```
blog df.columns
```

```
Out[15]:
```

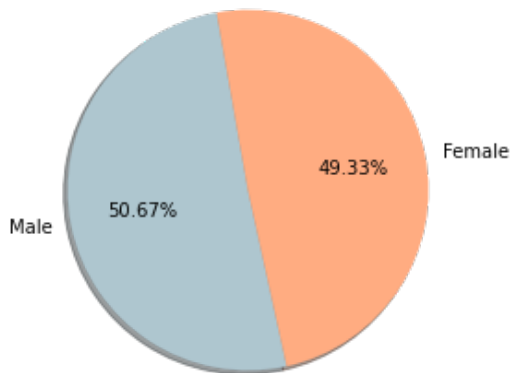
```
Index(['id', 'gender', 'age', 'topic', 'sign', 'date', 'text'], dtype='object')
```

```
In [16]:
```

```
# plotting the gender distribution
```

```
males = len(blog_df[blog_df['gender'] == 'male'])
females = len(blog_df[blog_df['gender'] == 'female'])
plt.pie(x=[males, females], explode=(0, 0), labels=['Male', 'Female'], autopct='%1.2f%%',
        shadow=True, startangle=100, colors=['#aec6cf', '#ffac81'])
fig = plt.gcf()
fig.set_size_inches(4.5, 4.5)
plt.title('Gender Distribution')
plt.show()
```

Gender Distribution

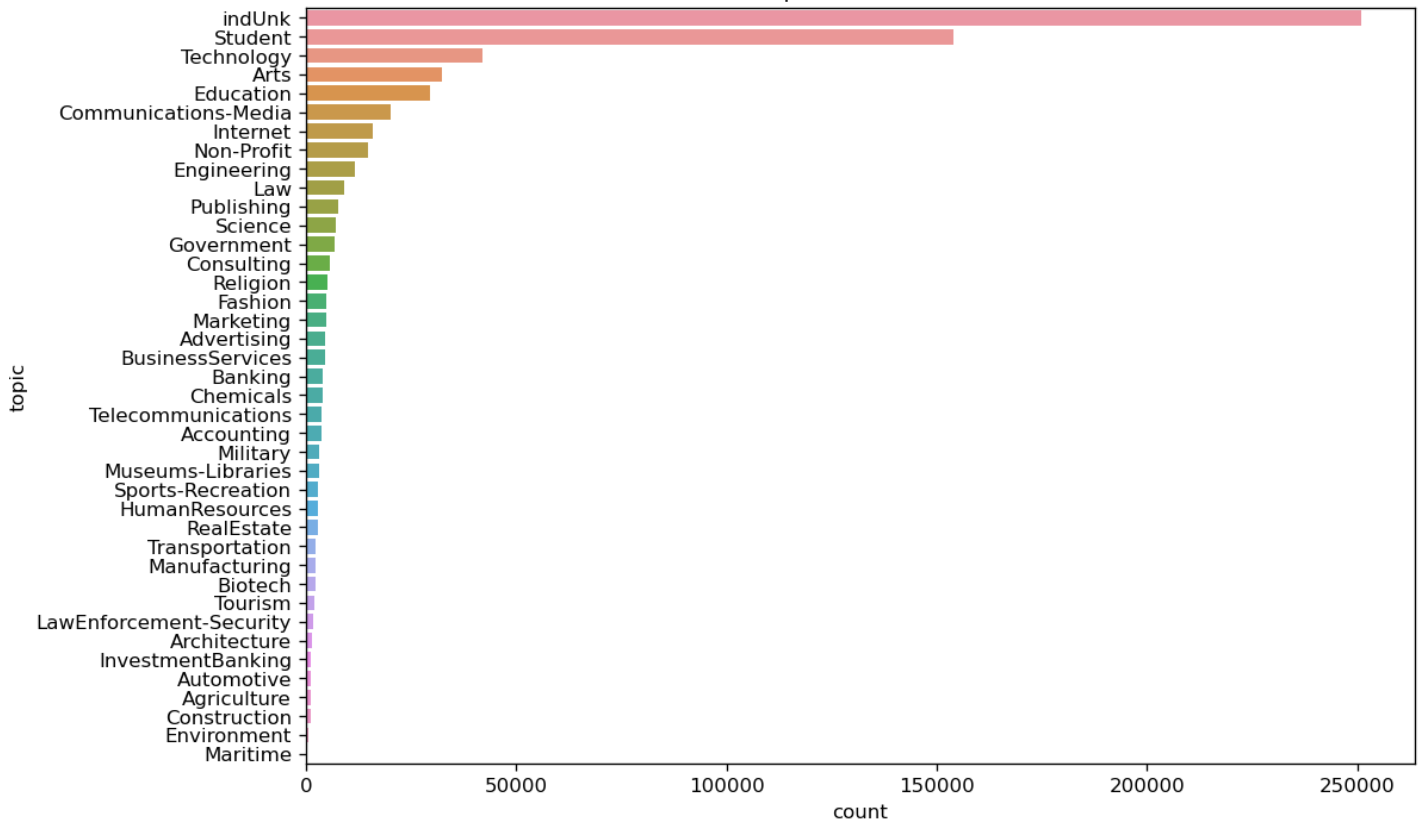


```
In [17]:
```

```
# plotting topic counts distribution
```

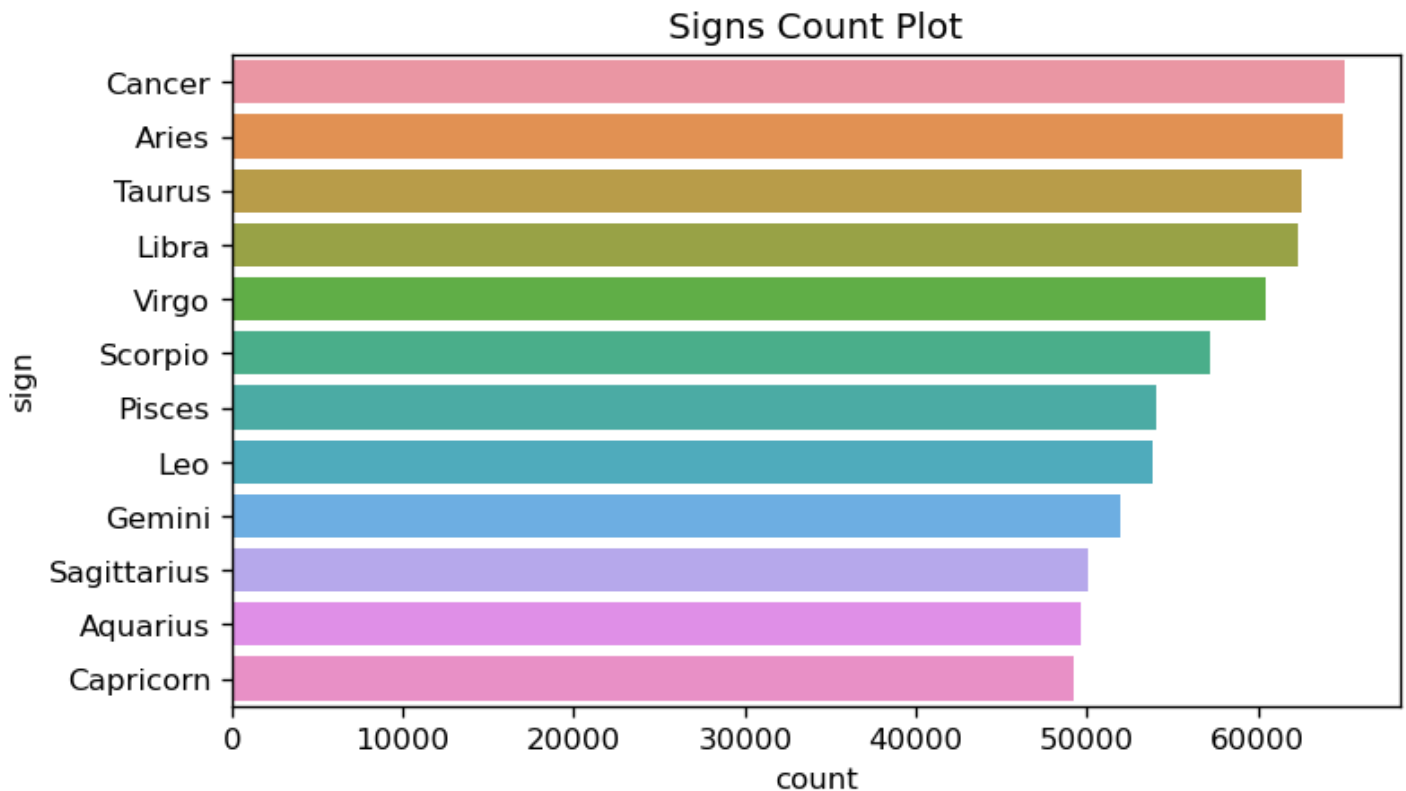
```
plt.figure(figsize=(10, 7), dpi=120)
sns.countplot(y='topic', data=blog_df, order=blog_df.topic.value_counts().index)
plt.title('Topics Count Plot')
plt.show()
```

Topics Count Plot



In [18]:

```
# plotting sign counts distribution
plt.figure(figsize=(7, 4), dpi=120)
sns.countplot(y='sign', data=blog_df, order=blog_df.sign.value_counts().index)
plt.title('Signs Count Plot')
plt.show()
```



In [19]:

```
from collections import Counter

# top 50 most frequent words in text
top_N = 50

words = (blog_df.text.str.cat(sep=' ').split())
rslt = pd.DataFrame(Counter(words).most_common(top_N),
                    columns=['Word', 'Frequency']).set_index('Word')
```

In [20]:

```
rslt[:50].transpose()
```

Out[20]:

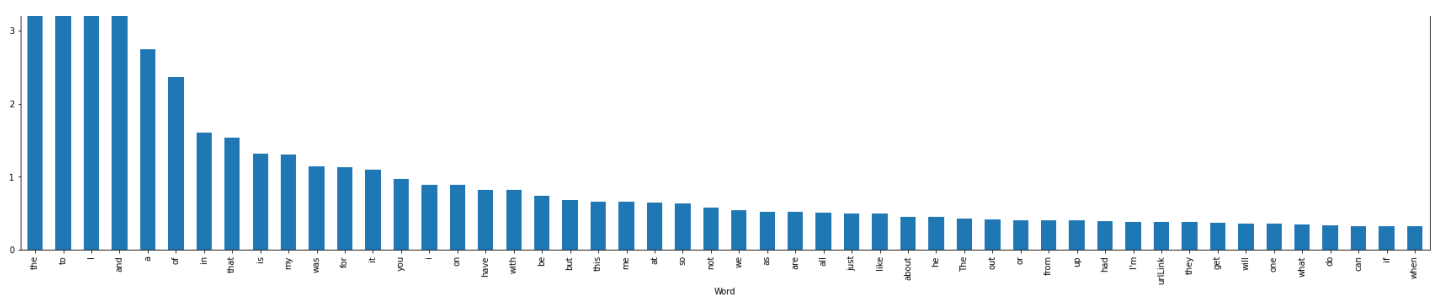
Word	the	to	I	and	a	of	in	that	is	my	...	urlLink	they
Frequency	4785809	3842654	3546837	3303447	2750057	2362870	1600803	1540799	1318972	1300183	...	383203	375996

1 rows × 50 columns

In [21]:

```
rslt.plot.bar(rot=0, figsize=(30,8), width=0.55)
plt.xticks(rotation=90)
plt.show()
```





In [22]:

```
pprint(rslt.index.tolist(), compact=True)
```

```
['the', 'to', 'I', 'and', 'a', 'of', 'in', 'that', 'is', 'my', 'was', 'for',
 'it', 'you', 'i', 'on', 'have', 'with', 'be', 'but', 'this', 'me', 'at', 'so',
 'not', 'we', 'as', 'are', 'all', 'just', 'like', 'about', 'he', 'The', 'out',
 'or', 'from', 'up', 'had', "I'm", 'urlLink', 'they', 'get', 'will', 'one',
 'what', 'do', 'can', 'if', 'when']
```

- **Many stopwords are occurring most frequently in the dataset, We will clean these in the preprocessing step**

In [23]:

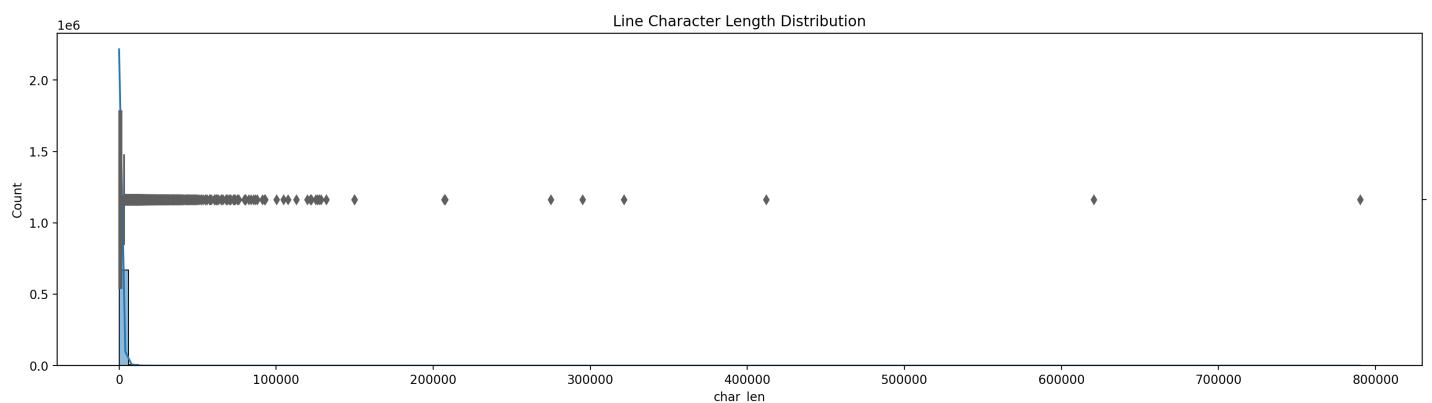
```
def get_text_len(row):
    row['char_len'] = len(row.text)
    row['word_len'] = len(row.text.split())
    return row
```

```
blog_df = blog_df.progress_apply(get_text_len, axis=1)
```

```
100%|████████████████████████████████████████████████████████████████████████████████| 681284/681
284 [11:02<00:00, 1028.75it/s]
```

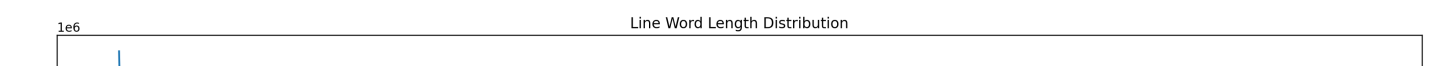
In [24]:

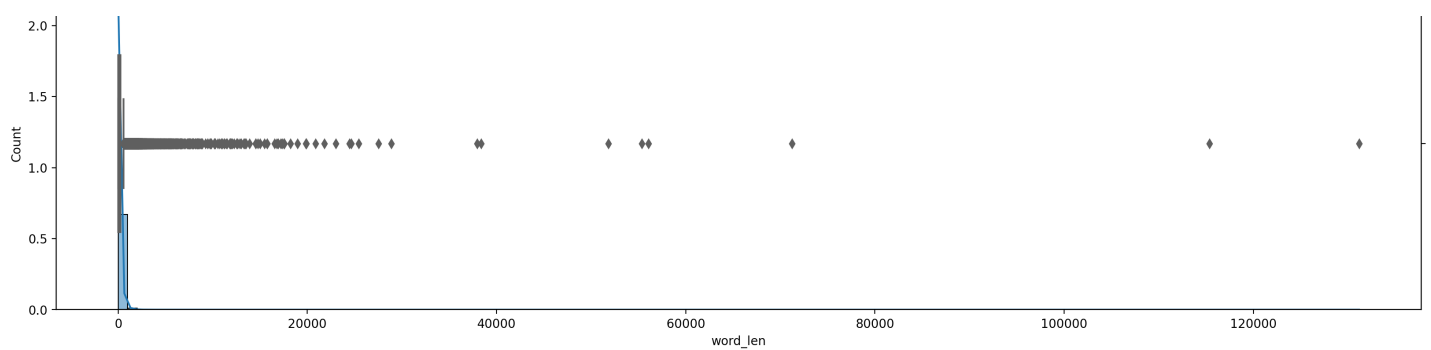
```
plt.figure(figsize=(20, 5), dpi=200)
ax = sns.histplot(x=blog_df.char_len, kde=True, bins=130)
ax2 = ax.twinx()
sns.boxplot(x=blog_df.char_len, ax=ax2, color='#EE6050')
ax2.set(ylim=(-.75, .75))
plt.title('Line Character Length Distribution')
plt.show()
```



In [25]:

```
plt.figure(figsize=(20, 5), dpi=200)
ax = sns.histplot(x=blog_df.word_len, kde=True, bins=130)
ax2 = ax.twinx()
sns.boxplot(x=blog_df.word_len, ax=ax2, color='#EE6050')
ax2.set(ylim=(-.75, .75))
plt.title('Line Word Length Distribution')
plt.show()
```





In [26]:

```
blog_df.describe()
```

Out[26]:

	id	age	char_len	word_len
count	6.812840e+05	681284.000000	681284.000000	681284.000000
mean	2.397802e+06	23.932326	1120.730698	200.786742
std	1.247723e+06	7.786009	2328.437003	415.160622
min	5.114000e+03	13.000000	4.000000	0.000000
25%	1.239610e+06	17.000000	230.000000	37.000000
50%	2.607577e+06	24.000000	637.000000	112.000000
75%	3.525660e+06	26.000000	1407.000000	255.000000
max	4.337650e+06	48.000000	790123.000000	131169.000000

- **Most blog texts have between 0 and 500 with median at 112 with relatively few outliers ranging till 1,31,169 words!**
This imbalance in text lengths will cause problem to classify the texts and we'll have to deal with it in preprocessing.

• Data Preprocessing

- **Data cleansing by removing unwanted characters, spaces, stop words etc. Convert text to lowercase**

In [27]:

```
# !pip install contractions
# !pip install beautifulsoup4
```

In [28]:

```
import nltk
nltk.download('punkt')
nltk.download('stopwords')

[nltk_data] Downloading package punkt to
[nltk_data] C:\Users\surya\AppData\Roaming\nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\surya\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

Out[28]:

True

In [29]:

```
# utility functions for text preprocessing
```

```

import re
import string
import unicodedata
import contractions
from bs4 import BeautifulSoup
from nltk.tokenize import word_tokenize, sent_tokenize
from nltk.stem.snowball import SnowballStemmer

CUSTOM = True

stemmer = SnowballStemmer('english')
if CUSTOM:
    stop_words = set(nltk.corpus.stopwords.words('english'))
    # custom stopwords added from the most frequent words which are generic
    # and might not relate to the sentiment of the review
    stop_words.update(['urllink'])
else:
    stop_words = set(nltk.corpus.stopwords.words('english'))

def replace_accented_chars(review_text):
    '''normalizes and replaces accented characters'''
    unaccented_text = unicodedata.normalize('NFKD', review_text).encode('ascii', 'ignore')
    unaccented_text = unaccented_text.decode('utf-8', 'ignore')
    return unaccented_text

def strip_html_tags(review_text):
    '''strips html tags like <h4> ..etc'''
    soup = BeautifulSoup(review_text, "html.parser")
    [s.extract() for s in soup(['iframe', 'script'])]
    stripped_text = soup.get_text()
    stripped_text = re.sub(r'[\r|\n|\r\n|+]', '\n', stripped_text)
    return stripped_text

def expand_contractions(review_text):
    review_text = contractions.fix(review_text)
    return review_text

def remove_special_characters(review_text):
    '''
    Remove special characters but preserve digits and exclamation marks
    as they indicate emotionally charged review '''
    review_text = re.sub(r"[^A-Za-z0-9!\?\\'\"]", " ", review_text)
    return review_text

def strip_stops(text, is_lower_case=False, stop_words=stop_words):
    '''strip stopwords'''
    tokens = word_tokenize(text)
    tokens = [token.strip() for token in tokens]
    if is_lower_case:
        filtered_tokens = [token for token in tokens if token not in stop_words]
    else:
        filtered_tokens = [token for token in tokens if token.lower() not in stop_words]
    filtered_text = ' '.join(filtered_tokens)
    return filtered_text

def snowball_stem(text, stemmer=stemmer):
    '''stemming using snowball stemmer'''
    words = text.split()
    stemmed_words = [stemmer.stem(word) for word in words]
    review_text = " ".join(stemmed_words)
    return review_text

```

In [30]:

```
def preprocess_text(text: str, lower=True, strip_punctuation=True) -> str:
```

	id	gender	age	topic	sign	date	text	char_len	word_len	clean_text
114792	3543301	female	24	Non-Profit	Gemini	03,August,2004	urlLink for real . i need to go h...	155	27	re: lun ma
599440	3535546	male	14	indUnk	Sagittarius	10,June,2004	634	80	so r r
241859	529513	male	33	Internet	Taurus	16,November,2003	This was the 4th or 5th time in my ...	122	17	4: ww m
321892	963380	male	24	Student	Cancer	20,January,2003	Good morning sunshine, as it's put in t...	1001	188	Q s pa dri

551120	1125546	id	gender	age	topic	sign	date	I must say, Barack Obama was impressive...	char_len	word_len	clear
30505	3687738		male	47	indUnk	Aries	12,agosto,2004	¿Acaso, el escenario descrito en el D...	648	103	des
552488	3056329		male	17	Government	Libra	15,May,2004	hey!! I havent blogged in awhile!...	704	125	r a t
380569	1417798		female	35	indUnk	Scorpio	15,September,2003	Hottie Alert! Greg K. just g...	698	119	 g
61256	2871824		male	24	Engineering	Leo	04,May,2004	Last 3 days quite happy lor because...	970	190	qu fa
150862	2427534		male	25	indUnk	Leo	02,July,2004	Where's my poker? Just finishe...	1143	176	po w
492401	3736084		male	25	Non-Profit	Leo	25,June,2004	Re: Laying to rest the idea of a pre-trib '...	1507	294	la pre '
127139	3541168		male	34	Government	Gemini	06,July,2004	and now has it's own blog at Zoe...	351	65	i p d
32151	1538911		female	35	indUnk	Libra	04,June,2004	urlLink All the mercy is right here...	228	32	
242087	2790498		female	23	indUnk	Aries	12,February,2004	I am moderately hung over. And ...	2044	381	m rain fri
616582	4296124		male	17	indUnk	Gemini	21,August,2004	urlLink Look at my boredom tonight!...	77	8	loo to hey
110967	3798616		female	26	indUnk	Pisces	29,June,2004	hey , just seeing if this works so i k...	855	176	v
586904	944569		female	14	Student	Leo	09,February,2004	Eeeexcellent...more Prince of Tenni...	890	140	 epi mi
624468	3653978		female	16	indUnk	Virgo	01,July,2004	I dunno y... whenever Im bullied by...	3589	677	wl ol re
388427	3168970		female	25	HumanResources	Cancer	19,July,2004	been slacking a bit the last few days...	4303	837	sk u t
456911	798653		male	26	indUnk	Aquarius	09,July,2004	You suck.	20	2	

```
top_N = 50
```

```
# top 50 most frequent words in cleaned text
words = (blog_df.cleaned_text.str.cat(sep=' ').split())
rslt = pd.DataFrame(Counter(words).most_common(top_N),
                    columns=['Word', 'Frequency']).set_index('Word')
```

In [35]:

```
rslt[:50].transpose()
```

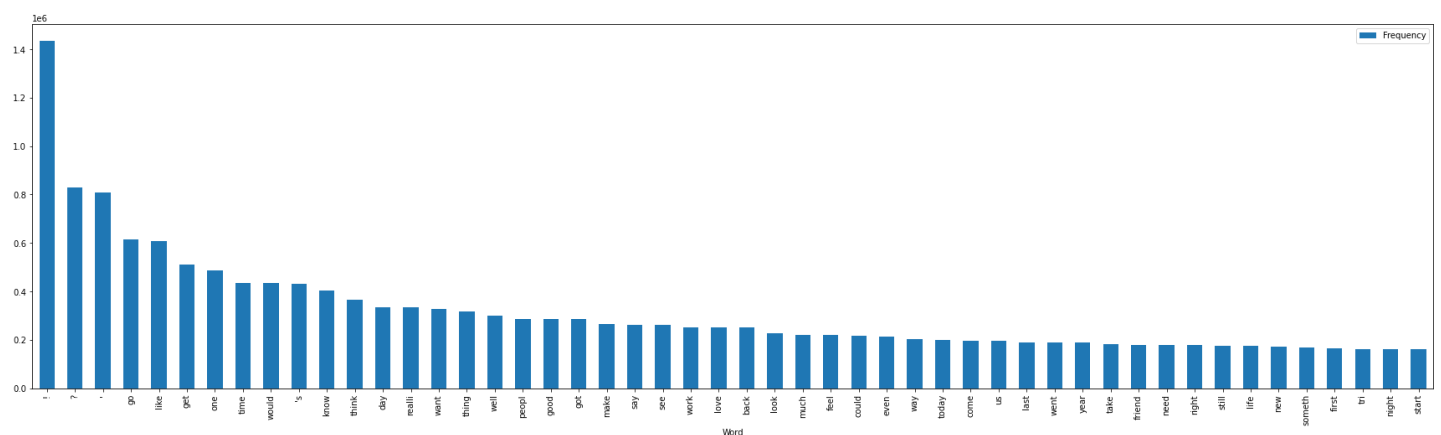
Out[35]:

	Word	!	?	'	go	like	get	one	time	would	's	...	need	right	still
Frequency	1433386	827744	807814	615740	606031	511351	487906	435766	433882	431960	...	176987	176689	175512	

1 rows × 50 columns

In [36]:

```
rslt.plot.bar(rot=0, figsize=(30,8), width=0.55)
plt.xticks(rotation=90)
plt.show()
```



In [37]:

```
pprint(rslt.index.tolist(), compact=True)
```

```
[['!', '?', '"', 'go', 'like', 'get', 'one', 'time', 'would', "'s", 'know',  
'think', 'day', 'realli', 'want', 'thing', 'well', 'peopl', 'good', 'got',  
'make', 'say', 'see', 'work', 'love', 'back', 'look', 'much', 'feel', 'could',  
'even', 'way', 'today', 'come', 'us', 'last', 'went', 'year', 'take', 'friend',  
'need', 'right', 'still', 'life', 'new', 'someth', 'first', 'tri', 'night',  
'start']]
```

- **Target/label merger and transformation**

In [80]:

```
# merge all labels together as we want to do multi-label classification
blog_df['labels'] = blog_df[['gender', 'age', 'topic', 'sign']].values.tolist()
```

In [81]:

```
dataset = blog_df.drop(columns = ['id', 'gender', 'age', 'topic', 'sign', 'date', 'text'])
```

- **Split dataset into train-test cuts**

In [82]:

```
from sklearn.model_selection import train_test_split
```

```
# Train-Test split of 80-20
```

```
X_train, X_test, y_train, y_test = train_test_split(dataset['cleaned_text'], dataset['labels'], test_size=0.20)
```

```
In [83]:
```

```
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
Out[83]:
```

```
((545027,), (136257,), (545027,), (136257,))
```

- **Vectorisation**

```
In [113]:
```

```
NGRAM_RANGE = (1, 2)
```

```
TOP_K = 30000
```

```
TOKEN_MODE = 'word'
```

```
MIN_DOC_FREQ = 2
```

```
kwargs = {  
    'ngram_range' : NGRAM_RANGE,  
    'dtype' : 'int32',  
    'strip_accents' : 'unicode',  
    'decode_error' : 'replace',  
    'analyzer' : TOKEN_MODE,  
    'min_df' : MIN_DOC_FREQ,  
    'max_features' : TOP_K  
}
```

```
In [114]:
```

```
# vectorize the texts to get features
```

```
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
from sklearn.feature_selection import SelectKBest, f_classif
```

```
vectorizer = TfidfVectorizer(**kwargs)
```

```
X_train_vec = vectorizer.fit_transform(X_train)
```

```
In [115]:
```

```
X_train_vec.shape
```

```
Out[115]:
```

```
(545027, 30000)
```

```
In [116]:
```

```
# do not fit on X_test to avoid data leakakge, only transform
```

```
X_test_vec = vectorizer.transform(X_test)
```

```
In [117]:
```

```
X_train_vec[0].shape
```

```
Out[117]:
```

```
(1, 30000)
```

```
In [118]:
```

```
X_test_vec[0].shape
```

```
Out[118]:
```

```
(1, 30000)
```


- **Transform labels**

In [119]:

```
y_train[0]
```

Out[119]:

```
['female', '23', 'Advertising', 'Taurus']
```

In [120]:

```
# use MultiLabelBinarizer to transform labels in a binary form so that the prediction will be a mask of 0s and 1s
from sklearn.preprocessing import MultiLabelBinarizer

binarizer = MultiLabelBinarizer()

# convert all labels to str
y_train = [[str(i) for i in j] for j in y_train]
y_test = [[str(i) for i in j] for j in y_test]

y_train_labels = binarizer.fit_transform(y_train)

# only transform test data to avoid data leakage
y_test_labels = binarizer.transform(y_test)

y_train_labels.shape, y_test_labels.shape
```

Out[120]:

```
((545027, 80), (136257, 80))
```

In [121]:

```
# converted to one hot vectors, each category here is a combination of labels from possible combinations of labels
y_train_labels[0]
```

Out[121]:

```
array([[0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0]])
```

- **Train Text Classifiers**

In [122]:

```
from sklearn.multiclass import OneVsRestClassifier
from sklearn.linear_model import LogisticRegression

from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score, classification_report
```

- **Logistic Regression**

In [123]:

```
# wrapper for lr classifier to be able to predict among classes
lr = OneVsRestClassifier(LogisticRegression(solver='lbfgs'), n_jobs=8)
```

In [124]:

```
lr.fit(X_train_vec, y_train_labels)
```

```
Out[124]:
```

```
OneVsRestClassifier(estimator=LogisticRegression(), n_jobs=8)
```

```
In [127]:
```

```
pred = lr.predict(X_test_vec)
```

```
In [128]:
```

```
# classification report
print(classification_report(y_test_labels, pred))
```

	precision	recall	f1-score	support
0	0.80	0.04	0.07	2696
1	0.66	0.06	0.11	5397
2	0.59	0.06	0.12	8288
3	0.68	0.11	0.19	14565
4	0.66	0.11	0.18	16374
5	0.69	0.03	0.05	14517
6	0.81	0.04	0.07	16103
7	0.69	0.02	0.05	13237
8	0.76	0.03	0.05	11033
9	0.74	0.03	0.05	9232
10	0.98	0.03	0.05	3456
11	0.90	0.11	0.20	4373
12	0.75	0.05	0.10	3481
13	0.87	0.07	0.12	2930
14	0.98	0.07	0.13	1835
15	0.91	0.03	0.06	1472
16	0.92	0.01	0.02	1072
17	0.98	0.22	0.36	1000
18	1.00	0.00	0.01	741
19	0.58	0.02	0.04	596
20	0.93	0.03	0.06	831
21	1.00	0.01	0.02	390
22	0.63	0.02	0.04	914
23	0.42	0.06	0.11	527
24	0.88	0.10	0.18	455
25	0.72	0.06	0.11	742
26	0.44	0.04	0.07	811
27	0.83	0.01	0.01	931
28	0.00	0.00	0.00	249
29	0.84	0.03	0.06	9879
30	1.00	0.00	0.01	347
31	0.76	0.04	0.07	13024
32	0.82	0.03	0.06	6467
33	0.00	0.00	0.00	246
34	1.00	0.00	0.01	830
35	1.00	0.03	0.05	437
36	0.88	0.05	0.09	905
37	0.86	0.06	0.10	12959
38	0.56	0.01	0.01	9755
39	0.00	0.00	0.00	751
40	0.71	0.01	0.02	3964
41	1.00	0.03	0.06	222
42	0.00	0.00	0.00	1195
43	0.81	0.04	0.07	5970
44	0.87	0.03	0.05	2325
45	0.00	0.00	0.00	119
46	0.97	0.16	0.27	966
47	0.77	0.02	0.04	10503
48	0.95	0.03	0.05	1422
49	0.50	0.00	0.00	627
50	0.81	0.02	0.05	3302
51	0.00	0.00	0.00	240
52	0.62	0.02	0.04	1782
53	1.00	0.00	0.01	385
54	0.90	0.02	0.04	10820
55	0.66	0.02	0.03	12313
56	1.00	0.00	0.01	442

56	1.00	0.00	0.01	442
57	0.00	0.00	0.00	48
58	0.64	0.01	0.01	979
59	1.00	0.00	0.01	601
60	0.56	0.02	0.04	600
61	0.55	0.03	0.05	3007
62	0.79	0.04	0.08	10768
63	0.92	0.07	0.14	1546
64	0.29	0.01	0.01	586
65	0.40	0.02	0.03	1084
66	0.69	0.02	0.03	10187
67	0.50	0.00	0.01	1406
68	0.76	0.02	0.04	11389
69	0.83	0.01	0.02	583
70	0.59	0.20	0.30	30722
71	0.87	0.03	0.05	12484
72	0.64	0.05	0.09	8401
73	0.87	0.03	0.05	800
74	1.00	0.02	0.04	377
75	1.00	0.13	0.23	457
76	0.83	0.03	0.05	12176
77	0.70	0.68	0.69	67216
78	0.55	0.21	0.30	50125
79	0.70	0.72	0.71	69041
micro avg	0.68	0.23	0.34	545028
macro avg	0.71	0.06	0.09	545028
weighted avg	0.71	0.23	0.26	545028
samples avg	0.68	0.23	0.33	545028

• Metrics:

Metrics for the model	Precision	Recall
Micro Average	0.68	0.23
Macro Average	0.71	0.06
Weighted Average	0.71	0.23

- **Micro-averaged Precision** is calculated as precision of Total values:
- all samples equally contribute to the final averaged metric
- **Macro-averaged Precision** is calculated as an average of Precisions of all classes:
- all classes equally contribute to the final averaged metric
- **Weighted-averaged Precision** is also calculated based on Precision per class but takes into account the number of samples of each class in the data:
- each classes's contribution to the average is weighted by its size

Which metric is relevant depends on If there is a class-imbalanced dataset? Is one class more important to get right than others? If you have an under-represented class which is important to your problem, macro-averaging may better, as it will highlight the performance of a model on all classes equally. On the other hand, if the assumption that all classes are equally important is not true, macro-averaging will over-emphasize the low performance on an infrequent class. Micro-averaging may be preferred in multilabel settings, including multiclass classification where a majority class is to be ignored.

- So, our model acheived a micro-avg f1-score of 0.34, macro-avg f1-score of 0.09. If we care about the minoority classes as well in the final classification, our model isn't up to the mark and might be further improved by cleaning the texts with multiple languages and performing oversampling/undersampling or SMOTE to deal with the class imbalance. Further, we can use BERT or other Attention-based classifiers which are much better at text classification tasks.

- **Print the true vs predicted labels for any 5 entries from the dataset**

In [140]:

```
from random import sample

# Getting real labels from transformed predicted labels
pred_classes = binarizer.inverse_transform(pred)
```

#Picking 5 random records from y_test and comparing actual labels vs predicted labels for those 5 records

```
for i in sample(range(len(pred)), 5):  
    print(i)  
    print("Actual labels: ", y_test[i])  
    print("Predicted labels:", pred_classes[i])  
    print()
```

78708

Actual labels: ['female', '14', 'Student', 'Virgo']

Predicted labels: ('female',)

129791

Actual labels: ['male', '27', 'Technology', 'Libra']

Predicted labels: ('male',)

90040

Actual labels: ['male', '24', 'Technology', 'Libra']

Predicted labels: ('male',)

117659

Actual labels: ['male', '24', 'Education', 'Leo']

Predicted labels: ('female', 'indUnk')

75481

Actual labels: ['female', '33', 'Communications-Media', 'Scorpio']

Predicted labels: ('male',)